

Darshan Tawari

SENIOR DATA ENGINEER | SPARK • SQL • PYTHON • AWS • GCP • GENAI • LLM

Phone: +91-7721905034 | Email: dtawari31@gmail.com | LinkedIn: [linkedin.com/in/dtawari31/](https://www.linkedin.com/in/dtawari31/)

Professional Summary

Senior Data Engineer with 7+ years of experience building scalable data platforms, distributed processing systems, and cloud-native analytics solutions on AWS and GCP. Skilled in SQL, Snowflake, Spark, Airflow, Python, Databricks, Data Modelling, ELT/ETL pipelines, and modern data platform architectures. Proven expertise in developing high-performance batch and real-time data pipelines, AI-powered data products, and enterprise analytics platforms processing millions of records daily. Hands-on experience with Generative AI, LLM applications, RAG architectures, and intelligent agent development. Strong background in data governance, data quality, platform automation, and partnering with product, engineering, and business stakeholders to deliver impactful data-driven solutions across multiple industry sectors, including Insurance, Retail, Sales, and Pharmaceutical.

Technical Skills

PROGRAMMING:

- Python, SQL, PySpark, Scala, Bash

BIG DATA & DISTRIBUTED SYSTEMS:

- Apache Spark, Apache Beam, Apache Airflow, Apache Kafka
- **Data Platforms & ETL** - Databricks, DBT, Snowflake, Datastage, Snaplogic.
- **Data Lake Formats** - Delta Lake, Apache Iceberg.

CLOUD PLATFORMS:

- **Amazon Web Service (AWS)** - Glue, EMR, API Gateway, Lambda, Step Function, CloudFormation, DMS, RDS, Redshift, Athena, Kinesis, DynamoDB, DataZone.
- **Google Cloud Platform (GCP)** - BigQuery, Cloud Composer, Cloud Function, Dataproc, Dataflow, PubSub, Vertex AI
- **IICS** - Informatica Intelligent Cloud Services.

GENAI & LLMs:

- OpenAI, Azure ChatGPT, LangChain, LLM integration, AI Agents, Salesforce Agentforce

Work Experience

Priorise

SENIOR DATA ENGINEER

Oct 2025 - Present

- Developing a **Salesforce Agentforce**-based agent in Salesforce to automate targeted Slack message delivery, using geographical and segment filters for audience selection, and leveraging **NLP** to refine audience targeting based on user-provided prompts.
- Developed an AI-powered **Agent** to recommend Next Best Actions (NBA) for sellers using LLMs and business intelligence data, identifying optimal actions based on open pipeline, cross-sell, renewal, and up-sell opportunities while leveraging ACV, quota attainment, and other key performance indicators (KPIs).
- Developing a scalable SQL-based Cohort Management platform to automate role assignment for newly onboarded sellers, ensuring accurate data processing, workflow automation, and seamless integration with downstream business systems.

Pixeldust Technologies

SENIOR DATA ENGINEER

Aug 2024 - Sep 2025

- Designed and implemented a real-time data pipeline for the **Telecom Domain** to process 100–150 million daily streaming events from various VoLTE network elements using **Pub/Sub**. Utilised **Apache Beam** on **Dataflow** with Flex Templates to aggregate performance counters in fixed 5-minute windows, enabling scalable and reusable pipeline deployments.
- Persisted the aggregated data in **BigQuery** and built a golden layer by applying domain-specific business logic. This refined dataset powers **Grafana** dashboards, which are used by operations teams for real-time visibility into critical network KPIs.
- Implemented automated anomaly detection with **SMTP** based email alerts to notify stakeholders of critical issues (e.g., call drops, registration failures), while simultaneously publishing alert messages to a **Kafka** topic for downstream processing and traceability.
- Implemented data quality checks, monitoring, validation frameworks, and automated alerting mechanisms to ensure reliability of business-critical datasets.

ZS Associates

SENIOR DATA ENGINEER

Nov 2021 - Jul 2024

- **Gen-AI:** Built a scalable Gen-AI platform for pre-clinical trial data, enabling document-level **Q&A chatbot** and **Text Summarization**. Developed robust pipelines for content extraction, chunking, and digitization, with embeddings stored in vector database. Integrated **Azure OpenAI** to support real-time document uploads, semantic search across uploaded and existing files, and generation of intelligent **LLM**-powered responses.
- **Data Mesh Architecture:** Spearheaded the implementation of a Data Mesh Architecture project using **AWS DataZone**, enabling domain teams to conduct cross-domain data analysis. Leveraged **Apache Iceberg** and AWS services including Glue, DMS, Step Functions, CloudFormation, Athena, and PySpark.
- **Platform Solutioning Project:** Served as the **Technical Lead** in designing a **custom web-app tool** for revenue forecasting and other metrics. Conducted POCs to evaluate requirements and technical feasibility, finalizing the design. Collaborated with business, UI, back-end, and data architect teams to create a **reusable, cloud-based application** that reduced operational expenses by **33%**.
- **Infrastructure Automation:** **Led the team** to automate infrastructure setup for new projects within the client's environment using **Python** and **APIs**. This automation included creating **IICS** project folders and mappings, setting up **Tidal** scheduler jobs, and provisioning **AWS services** such as S3 buckets, EC2 clusters, RDS databases, DynamoDB databases, and other services as needed.
- **LangChain LLM:** Led an in-house POC for LangChain Large Language Model (LLM) using **GPT-4** and **AzureChatOpenAI**. Utilized SqlDatabase-Toolkit with PythonAgentExecutor and SqlAgentExecutor to generate SQL queries or Python code for visualization charts from text-based inputs. Developed a **Streamlit** application to facilitate a chat-like interaction with the model.
- **Cloud-Native Architecture:** Designed, developed, and deployed a cost-effective, cloud-native architecture with centralized data governance. Utilized AWS API Gateway, Lambda, S3, CloudFront, and Route 53 to streamline operations.

Lumiq.AI

DATA ENGINEER

Nov 2020 - Nov 2021

- **PySpark Framework Development:** Developed a PySpark framework for migrating incremental data from AWS S3 to **AWS Kinesis** streams. This framework triggered AWS **Lambda** functions to transform data based on business logic and store it in **DynamoDB**.
- **AWS CloudWatch Dashboard:** Designed an AWS CloudWatch dashboard to monitor **AWS Lambda** function errors, invocations, exceptions, and **AWS DMS** metrics, providing insights into data latency for each table.
- **Testing Environment Setup:** Created a testing environment leveraging **AWS CLI** by the generic scripts, enabling the replication of existing AWS services such as Lambda functions, Kinesis streams, Glue jobs, and DynamoDB tables on UAT, dynamic tags, and primary and secondary indexes for DynamoDB tables.
- **Enterprise Data Lake Creation:** Established an enterprise data lake by migrating data from legacy on-premise servers to the AWS cloud platform. Developed data ingestion processes using AWS DMS from multiple sources including Oracle, MySQL, and SQL Server. Implemented business logic to write efficient Spark jobs for data transformation using Amazon Glue.

Datametica Birds

BIGDATA ENGINEER

Jul 2019 - Nov 2020

- Converted the **Datastage** jobs to facilitate the loading and transformation of data into BigQuery, with scheduling managed through **Cloud Composer**.
- Developed and deployed end-to-end pipelines for transforming, migrating, and loading data from **Teradata** into **BigQuery**. Implemented business logic and aggregations using BigQuery. Automated ETL pipelines using **Cloud Composer** and **IBM Control-M**.
- Developed a **Python framework** to recursively trigger multiple **DAGs** from a single DAG in Google Cloud Composer, addressing various project needs.
- **Bug Resolution and Hot-fix Team:** Worked in the Hot-fix team to resolve a significant number of client-raised bugs, contributing to project stability.

Education

Pune Institute of Computer Technology, Pune University (SPPU)

Pune, India

BACHELOR OF ENGINEERING IN INFORMATION TECHNOLOGY.

2015 - 2019

- Graduated with **First Class with Distinction**
- Scored **7.9 CGPA**

Certifications

- **GCP** : Professional Database Engineer, Professional Data Engineer
- **AWS** : Cloud Associate Developer, Cloud Practitioner